

```
In [205]: import gym
import pybulletgym
import pybulletgym.envs
import numpy as np
import math
import matplotlib.pyplot as plt
from numpy.linalg import pinv
import time
```

```
In [212]: import torch
import torch.nn as nn
import torch.optim as optim
import torch.nn.functional as F
from torch.autograd import Variable
from torch.distributions import Categorical
```

```
In [207]: env = gym.make("CartPole-v1")
env.reset()
```

```
Out[207]: array([-0.03294539,  0.04493735,  0.00958779,  0.01345614])
```

```

In [217]: learning_rate = 0.01
          gamma = 0.99

          #4 states, 2 actions
          num_states = env.observation_space.shape[0]
          num_actions = env.action_space.n

          class PolicyNetwork(nn.Module):
              def __init__(self):
                  super(PolicyNetwork, self).__init__()
                  self.state_num = num_states
                  self.action_num = num_actions

                  self.l1 = nn.Linear(self.state_num, 128, bias=False)
                  self.l2 = nn.Linear(128, self.action_num, bias=False)

                  self.gamma = gamma

                  # Episode policy and reward history
                  self.policy_episode = Variable(torch.Tensor())
                  self.reward_episode = [] #a hash containing the reward of step
                  # Overall return history
                  self.return_history = []
                  self.return_reward = []
                  self.return_history_stepzero = []
                  self.episodesPeriter = []

              def forward(self, x):
                  model = torch.nn.Sequential(
                      self.l1,
                      nn.Dropout(p=0.6),
                      #nn.ReLU(),
                      nn.Tanh(),
                      self.l2,
                      nn.Softmax(dim=-1)
                  )

                  return model(x)

```

In [218]: `class CartPole ():`

```

def __init__(self, env, policy, part, totsteps, iterationsNo, learning_rate, gamma):
    self.env = env
    self.policy = policy
    self.part = part
    self.learning_rate = learning_rate
    self.gamma = gamma
    self.optimizer = optim.Adam(self.policy.parameters(), lr=learning_rate)
    self.iterationsNo = iterationsNo
    self.totsteps = totsteps

    self.iterations = []

    self.polHist_all episodes = Variable(torch.Tensor())
    self.rewHist_all episodes = Variable(torch.Tensor())
    self.rewHist_all episodes_mod = Variable(torch.Tensor())

def select_action(self, state):
    #Select an action (0 or 1) by running policy model and choosing
    state = torch.from_numpy(state).type(torch.FloatTensor)
    act_prob = self.policy(state)
    c = Categorical(act_prob)
    action = c.sample()

    if len(self.policy.policy_episode) > 0:
        self.policy.policy_episode = torch.cat([self.policy.policy_episode, action])
    else:
        self.policy.policy_episode = (c.log_prob(action).reshape(1))
    return action

def rewardFunction(self, polHistory, rewHistory, part = 1):

    Rewards_tot = 0
    rewards = []
    for Reward in rewHistory[::-1]:
        Rewards_tot = Reward + self.policy.gamma * Rewards_tot
        rewards.insert(0, Rewards_tot)

    if(self.part == 1):
        #gainPerStep = rewards[0]
        gainPerStep = sum(rewHistory)
        gainZeroStep = rewards[0]
        reward = rewards[0]*torch.sum(polHistory)

    if(self.part == 2):
        gainPerStep = sum(rewHistory)
        gainZeroStep = rewards[0]
        rewards = torch.FloatTensor(rewards)
        reward = torch.sum(torch.mul(polHistory, rewards))

    if(self.part == 3):

```

```

gainPerStep = sum(rewHistory)
gainZeroStep = rewards[0]
rewards = torch.FloatTensor(rewards)
#rewards = (rewards - rewards.mean()) / (rewards.std())
#reward = torch.sum(torch.mul(polHistory, rewards))

#for i in range(0, len(polHistory)):
#print(len(self.polHist_all episodes))
#print(len(self.rewHist_all episodes))

#if self.polHist_all episodes.size(0) > 0:
if (len(self.polHist_all episodes) > 0):
    #print(self.polHist_all episodes.shape)
    #print(type(polHistory))
    self.polHist_all episodes = torch.cat([self.polHist_all episodes, polHistory])
    self.rewHist_all episodes = torch.cat([self.rewHist_all episodes, rewards])
else:
    self.polHist_all episodes = polHistory
    self.rewHist_all episodes = rewards
#rewards = (rewards - rewards.mean()) / (rewards.std())
#reward = torch.sum(torch.mul(polHistory, rewards))

#self.rewHist_all episodes = (self.rewHist_all episodes - self.rewHist_all episodes_mod)
self.rewHist_all episodes_mod = (self.rewHist_all episodes - self.rewHist_all episodes_mod)
reward = torch.sum(torch.mul(self.polHist_all episodes, self.rewHist_all episodes))
#print(reward)

return reward, gainPerStep, gainZeroStep

def update_policy(self, reward, retTraj_tot, retTraj_tot_stepzero, episodes_iter):
    #print("in update")
    # Update network weights
    self.optimizer.zero_grad()
    #print(reward)
    reward.backward()
    self.optimizer.step()
    self.policy.return_history.append(retTraj_tot)
    #self.policy.return_reward.append(reward)
    self.policy.return_history_stepzero.append(retTraj_tot_stepzero)
    self.policy.episodesPeriter.append(episodes_iter)

def reinforceAlgo(self):
    #running_reward = 10

    for iter in range(self.iterationsNo):
        self.polHist_all episodes = Variable(torch.Tensor())
        self.rewHist_all episodes = Variable(torch.Tensor())
        #print("iteration no", iter)
        steps = 0
        state = env.reset() # Reset environment, starting state record
        done = False
        episodes = 0;
        rewardFunc = Variable(torch.FloatTensor([0]))
        rewardEpisode = Variable(torch.FloatTensor())
        retTraj_tot = 0

```

```

retStepzero_tot = 0
#optimizer.zero_grad()
while(steps < self.totsteps):
    steps += 1;
    action = self.select_action(state)
    # Step through environment using chosen action
    state, reward, done, _ = env.step(action.item())
    #env.render()

# Save reward
self.policy.reward_episode.append(reward)
if (done == True):
    rewardEpisode, retTraj, retStepzero = self.rewardFunc(state, action, reward)
    if(self.part != 3):
        rewardFunc += rewardEpisode
        retTraj_tot += retTraj
        retStepzero_tot += retStepzero
    #reset the environment again
    self.policy.policy_episode = Variable(torch.Tensor([0]))
    self.policy.reward_episode = []
    state = env.reset()
    done = False
    episodes += 1

if(self.part == 3):
    rewardFunc += rewardEpisode
if(episodes > 0):
    self.update_policy(-1 * rewardFunc/episodes, retTraj_tot, retStepzero_tot)
else:
    self.update_policy(-1 * rewardFunc, retTraj_tot, retStepzero_tot)
self.iterations.append(iter)

```

#####1.1 - Batch size = 500, iterations = 200

#####

In [219]:



```

policy = PolicyNetwork()
cartPole = CartPole(env= env, policy = policy, part = 1, totsteps = 500)
cartPole.reinforceAlgo()

```

```
In [220]: plt.plot(cartPole.iterations, cartPole.policy.return_history_stepzero, color='b')
plt.xlabel("iterations")
plt.ylabel("avreturn_trajectory")
plt.show()

plt.plot(cartPole.iterations, cartPole.policy.episodesPeriter, color='b')
plt.xlabel("iterations")
plt.ylabel("episodes")
plt.show()
```



```
#####1.2 - Batch size = 500, iterations = 200  
#####
```



```

In [222]: policy = PolicyNetwork()
cartPole = CartPole(env= env, policy = policy, part = 2, totsteps = 500)
cartPole.reinforceAlgo()

plt.plot(cartPole.iterations, cartPole.policy.return_history_stepzero, color='b')
plt.xlabel("iterations")
plt.ylabel("avreturn_trajectory")
plt.show()

plt.plot(cartPole.iterations, cartPole.policy.episodesPeriter, color='b')
plt.xlabel("iterations")
plt.ylabel("episodes")
plt.show()

```



```

#####1.3 - Batch size = 500, iterations = 200
#####



```



```
In [224]: policy = PolicyNetwork()
cartPole = CartPole(env= env, policy = policy, part = 3, totsteps = 500)
cartPole.reinforceAlgo()

plt.plot(cartPole.iterations, cartPole.policy.return_history_stepzero, color='b')
plt.xlabel("iterations")
plt.ylabel("avreturn_trajectory")
plt.show()

plt.plot(cartPole.iterations, cartPole.policy.episodesPeriter, color='b')
plt.xlabel("iterations")
plt.ylabel("episodes")
plt.show()
```





```
#####1.3 - Batch size = 600, iterations = 200#####
```

```
In [225]: policy = PolicyNetwork()
cartPole = CartPole(env= env, policy = policy, part = 3, totsteps = 600)
cartPole.reinforceAlgo()

plt.plot(cartPole.iterations, cartPole.policy.return_history_stepzero, color='b')
plt.xlabel("iterations")
plt.ylabel("avreturn_trajectory")
plt.show()

plt.plot(cartPole.iterations, cartPole.policy.episodesPeriter, color='b')
plt.xlabel("iterations")
plt.ylabel("episodes")
plt.show()
```



```
#####1.3 - Batch size = 800, iterations = 200#####
```



```

In [228]: policy = PolicyNetwork()
cartPole = CartPole(env= env, policy = policy, part = 3, totsteps = 800)
cartPole.reinforceAlgo()

plt.plot(cartPole.iterations, cartPole.policy.return_history_stepzero, color='b')
plt.xlabel("iterations")
plt.ylabel("avreturn_trajectory")
plt.show()

plt.plot(cartPole.iterations, cartPole.policy.episodesPeriter, color='b')
plt.xlabel("iterations")
plt.ylabel("episodes")
plt.show()

```



```



#####1.3 - Batch size = 1000, iterations = 200
#####

```

```
In [230]: policy = PolicyNetwork()
cartPole = CartPole(env= env, policy = policy, part = 3, totsteps = 1000)
cartPole.reinforceAlgo()

plt.plot(cartPole.iterations, cartPole.policy.return_history_stepzero, color='b')
plt.xlabel("iterations")
plt.ylabel("avreturn_trajectory")
plt.show()

plt.plot(cartPole.iterations, cartPole.policy.episodesPeriter, color='b')
plt.xlabel("iterations")
plt.ylabel("episodes")
plt.show()
```



Increasing the batch size ideally should give better training results, since we would be exploring more no of trajectories in a given iteration to determine the trajectory with the best result. However, in this case the average discounted reward per episode remains 100 in all the above four cases, meaning the training has achieved the result with batch size 500 only.

If we look at rewards obtained at each step the average max reward in the case of 500 steps would be 500, while in the case of 600,800 and 1000 steps it would be 300,400 and 500 respectively since the min no of episodes would be more than 1. Hence that cannot be given as the metric to determine the performance of training. Hence the average discounted measure would tell us the training performance, which is equivalent in all the above cases.

```
In [1]: import gym
import pybulletgym
import pybulletgym.envs
import numpy as np
import math
import matplotlib.pyplot as plt
from numpy.linalg import pinv
import time
```

```
In [2]: import torch
import torch.nn as nn
import torch.optim as optim
import torch.nn.functional as F
from torch.autograd import Variable
from torch.distributions import Categorical
from torch.distributions import MultivariateNormal
```

```
In [3]: env = gym.make("modified_gym_env:ReacherPyBulletEnv-v1", rand_init = False)
#env.render()
env.reset()
```

```
current_dir=/home/apurba/.virtualenvs/276c_assgn/lib/python3.6/site-packages/pybullet_envs/bullet
options=
```

```
Out[3]: array([ 0.3928371 ,  0.3928371 , -0.68091764,  0.26561381,  0.5
,
0.          ,  0.08333333,  0.          ])
```

```
In [4]: learning_rate = 0.01
gamma = 0.90

num_states = env.observation_space.shape[0]
print(num_states)
```

9

```

In [5]: class PolicyNetwork(nn.Module):
        def __init__(self):
            super(PolicyNetwork, self).__init__()
            self.state_num = 8
            self.action_num = 2

            #self.l1 = nn.Linear(self.state_num, 64, bias=False)
            #self.l2 = nn.Linear(64, 64, bias=False)
            #self.l3 = nn.Linear(64, self.action_num, bias=False)
            self.l1 = nn.Linear(self.state_num, 128, bias=False)
            self.l2 = nn.Linear(128, self.action_num, bias=False)
            self.gamma = gamma

            # Episode policy and reward history
            self.policy_episode = Variable(torch.Tensor())
            self.reward_episode = []
            # Overall return history
            self.return_history = []
            self.return_reward = []
            self.return_history_stepzero = []

            self.policy_episode_eval = Variable(torch.Tensor())
            self.reward_episode_eval = []

            self.episodesPeriter = []
            self.sigmax = torch.nn.Parameter(torch.FloatTensor([0.1]))
            self.sigmay = torch.nn.Parameter(torch.FloatTensor([0.1]))
            #self.sigma = torch.nn.Parameter(torch.tensor([[0.1, 0], [0, 0.1]]))

        def forward(self, x):

            model = torch.nn.Sequential(
                self.l1,
                nn.Dropout(p=0.6),
                #nn.ReLU(),
                nn.Tanh(),
                self.l2,
                #nn.Dropout(p=0.6),
                #nn.Tanh(),
                #self.l3,
                nn.Tanh()
            )

            return model(x)

```

In [6]: `class modtwoLink ():`

```

def __init__(self, env, policy, part, totsteps, iterationsNo, learning_rate, gamma):
    self.env = env
    self.policy = policy
    self.part = part
    self.learning_rate = learning_rate
    self.gamma = gamma
    self.optimizer = optim.Adam(self.policy.parameters(), lr=learning_rate)
    self.iterationsNo = iterationsNo
    self.totsteps = totsteps

    self.iterations = []

    self.polHist_all episodes = Variable(torch.Tensor())
    self.rewHist_all episodes = Variable(torch.Tensor())
    self.rewHist_all episodes_mod = Variable(torch.Tensor())

def select_action(self, state):

    #Select an action (0 or 1) by running policy model and choosing
    state = torch.from_numpy(state).type(torch.FloatTensor)
    probs = self.policy(state)#this will give the mean of x and y
    mux = policy.sigmax.reshape(1)
    muy = policy.sigmay.reshape(1)
    if(mux < 0.001):
        mux = mux + 0.001
    if(muy < 0.001):
        muy = muy + 0.001
    covariance = torch.cat([mux, muy])
    #print(covariance)
    covariance_tensor = torch.FloatTensor(covariance)
    c = MultivariateNormal(probs, torch.diag(torch.abs(covariance_tensor)))
    #print(self.policy.sigma)
    #c = MultivariateNormal(probs, self.policy.sigma)
    action = c.sample()

    if len(self.policy.policy_episode) > 0:
        self.policy.policy_episode = torch.cat([self.policy.policy_episode, action])
    else:
        self.policy.policy_episode = (c.log_prob(action).reshape(1))
    return action

def rewardFunction(self, polHistory, rewHistory, part = 1):

    Rewards_tot = 0
    rewards = []
    for Reward in rewHistory[::-1]:
        Rewards_tot = Reward + self.policy.gamma * Rewards_tot
        rewards.insert(0, Rewards_tot)

    if(self.part == 1):
        #gainPerStep = rewards[0]

```



```

gainPerStep = sum(rewHistory)
gainZeroStep = rewards[0]
reward = rewards[0]*torch.sum(polHistory)

if(self.part == 2):
    gainPerStep = sum(rewHistory)
    gainZeroStep = rewards[0]
    rewards = torch.FloatTensor(rewards)
    reward = torch.sum(torch.mul(polHistory, rewards))

if(self.part == 3):
    gainPerStep = sum(rewHistory)
    gainZeroStep = rewards[0]
    rewards = torch.FloatTensor(rewards)

    if (len(self.polHist_allepisodes) > 0):

        self.polHist_allepisodes = torch.cat([self.polHist_allepisodes, polHistory])
        self.rewHist_allepisodes = torch.cat([self.rewHist_allepisodes, rewards])
    else:
        self.polHist_allepisodes = polHistory
        self.rewHist_allepisodes = rewards

    #self.rewHist_allepisodes = (self.rewHist_allepisodes - self.rewHist_allepisodes_mod)
    self.rewHist_allepisodes_mod = (self.rewHist_allepisodes - self.rewHist_allepisodes_mod)
    reward = torch.sum(torch.mul(self.polHist_allepisodes, self.rewHist_allepisodes))
    #print(reward)

return reward, gainPerStep, gainZeroStep

def update_policy(self, reward, retTraj_tot, retTraj_tot_stepzero, episodes_iter):
    #print("in update")
    # Update network weights
    self.optimizer.zero_grad()
    #print(reward)
    reward.backward()
    self.optimizer.step()
    self.policy.return_history.append(retTraj_tot)
    #self.policy.return_reward.append(reward)
    self.policy.return_history_stepzero.append(retTraj_tot_stepzero)
    self.policy.episodesPeriter.append(episodes_iter)

def reinforceAlgo(self):
    #running_reward = 10

    for iter in range(self.iterationsNo):
        self.polHist_allepisodes = Variable(torch.Tensor())
        self.rewHist_allepisodes = Variable(torch.Tensor())
        #print("iteration no",iter)
        steps = 0
        state = env.reset() # Reset environment, starting state record
        done = False

```

```

episodes = 0;
rewardFunc = Variable(torch.FloatTensor([0]))
rewardEpisode = Variable(torch.FloatTensor())
retTraj_tot = 0
retStepzero_tot = 0
#optimizer.zero_grad()
while(steps < self.totsteps):
    steps += 1;
    action = self.select_action(state)
    # Step through environment using chosen action
    state, reward, done, _ = env.step(action)
    #env.render()

# Save reward
self.policy.reward_episode.append(reward)
if (done == True):
    rewardEpisode, retTraj, retStepzero = self.rewardFunc
    if(self.part != 3):
        rewardFunc += rewardEpisode
        retTraj_tot += retTraj
        retStepzero_tot += retStepzero
    #reset the environment again
    self.policy.policy_episode = Variable(torch.Tensor(
    self.policy.reward_episode = []
    state = env.reset()
    done = False
    episodes += 1

if(self.part == 3):
    rewardFunc += rewardEpisode
if(episodes > 0):
    self.update_policy(-1 * rewardFunc/episodes, retTraj_tot
else:
    self.update_policy(-1 * rewardFunc, retTraj_tot, retStep
self.iterations.append(iter)

```



```

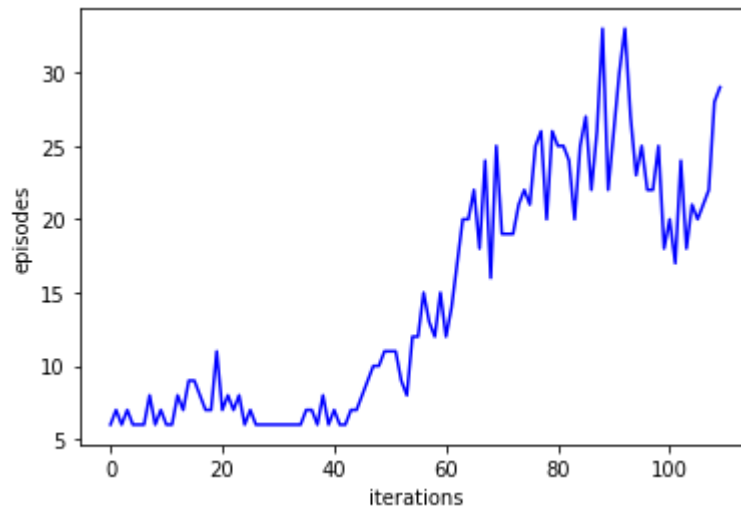
In [7]: state = env.reset()
policy = PolicyNetwork()
twoLink = modtwoLink(env= env, policy = policy, part = 3, totsteps = 1000)
twoLink.reinforceAlgo()

```

The below graphs denote the average rewards, average discounted reward and the no of episodes varying with the iteration. The no of episodes and rewards should increase with no of iterations.

```
In [8]: plt.plot(twoLink.iterations,twoLink.policy.return_history, color='b');  
plt.xlabel("iterations")  
plt.ylabel("return_trajectory")  
plt.show()  
  
plt.plot(twoLink.iterations,twoLink.policy.return_history_stepzero, color='b');  
plt.xlabel("iterations")  
plt.ylabel("return_trajectory")  
plt.show()  
  
plt.plot(twoLink.iterations,twoLink.policy.episodesPeriter, color='b');  
plt.xlabel("iterations")  
plt.ylabel("episodes")  
plt.show()
```





For evaluation, sampling the x, y values from the mean obtained from the network

```
In [9]: def select_action_eval(state):  
        #Select an action (0 or 1) by running policy model and choosing base  
        state = torch.from_numpy(state).type(torch.FloatTensor)  
        action = policy(state)#this will give the mean of x and y  
  
        return action
```

False

```
False  
False  
False  
False  
False  
False  
False  
False  
True  
49
```

```
In [ ]: env.close()
```